

Lecture 8

Saurabh

Abstract

This guide aims to demystify backpropagation and linearization, two foundational concepts in machine learning. We will move from high-level intuition to a concrete, step-by-step mathematical understanding. Our goal is to show not just *what* the formulas are, but *why* they are what they are, in a way that is clear, simple, and memorable for any student.

The Core Idea: How Do Machines Learn?

Imagine you're playing a game of darts. Your first throw is off the mark. To improve, your brain subconsciously performs a series of calculations. It sees *how far* off you were and in which direction (this is the **error**), and you adjust your stance, aim, angle, and force for the next throw. You are learning from your error.

In machine learning, this process is formalized and given mathematical rigor:

- **The Throw:** The model makes a prediction based on its current state.
- **The Error:** A **loss function** measures how far the prediction is from the actual target. A high loss means a bad prediction; a low loss means a good one.
- **The Adjustment:** The model must adjust its internal **parameters** (weights and biases) to reduce this error for the next prediction.

The key question is: *how exactly should it adjust them?* If we have millions of parameters, which ones do we change, and by how much?

This is where the **gradient** comes in. Think of the loss function as a hilly landscape, where your goal is to get to the lowest valley (minimum loss). The gradient at your current position is a vector that points in the direction of the *steepest ascent*—the fastest way to go uphill. To reduce the error, we simply take a small step in the **exact opposite direction** of the gradient.

This entire guide is about **Backpropagation**, the powerful and brilliantly efficient algorithm that calculates these necessary gradients for every single parameter in the model, telling us exactly how to "adjust our aim."

1. The Challenge: Differentiation is Hard

For simple functions, finding a derivative is easy. But what about a complex function composed of many nested parts, like the one from the lecture?

$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$

Manually deriving this using the standard chain rule gives a very long and messy expression. Implementing this by hand is not only slow but also extremely prone to errors. Now, consider a real-world neural network like GPT-3, which has *billions* of parameters. A manual approach is not just impractical; it's fundamentally impossible. We need an automated, efficient, and scalable method.

2. The Solution: Automatic Differentiation

Instead of tackling the whole function at once, we can break it down into a series of simple, elementary operations whose derivatives we already know. This creates a roadmap for the calculation, which we call a **computation graph**.

2.1 The Forward Pass: Building the Function Step-by-Step

Let's decompose our function. This process is called the **forward pass**, where we start with an input and compute the final value step-by-step. Let's make this concrete with a numerical example: let $x = 2$.

- Start with input $x = 2$
- Let $a = x^2 \rightarrow a = 2^2 = 4$
- Let $b = \exp(a) \rightarrow b = \exp(4) \approx 54.6$
- Let $c = a + b \rightarrow c = 4 + 54.6 = 58.6$
- Let $d = \sqrt{c} \rightarrow d = \sqrt{58.6} \approx 7.66$
- Let $e = \cos(c) \rightarrow e = \cos(58.6) \approx 0.83$
- Finally, $f = d + e \rightarrow f = 7.66 + 0.83 = 8.49$

This sequence is visualized in the computation graph below. The green values show the result of the forward pass.

2.2 The Backward Pass: The Magic of the Chain Rule

Now, we find $\frac{df}{dx}$ by moving **backward** through the graph, starting from the end. The core idea is simple: at each node, we compute its gradient, which represents that node's total influence on the final output f . We start with the derivative of the output with respect to itself, which is trivially 1 ($\frac{df}{df} = 1$).

Then, for any node u that is an input to a node v , the chain rule tells us:


$$\frac{\partial f}{\partial u} = \frac{\partial f}{\partial v} \cdot \frac{\partial v}{\partial u}$$

If a node u feeds into multiple other nodes (v_1, v_2, \dots), its total influence is the sum of its influences through each path (multivariate chain rule):

$$\frac{\partial f}{\partial u} = \sum_i \frac{\partial f}{\partial v_i} \frac{\partial v_i}{\partial u}$$

This reuse of previously computed "downstream" gradients ($\frac{\partial f}{\partial v_i}$) is what makes back-propagation so efficient.

Let's apply this to our example. We'll denote $\frac{df}{dy}$ as \bar{y} . We start with $\bar{f} = 1$. The **red values** on the graph show the gradients calculated during the backward pass.



Backward Pass Step-by-Step Calculation:

1. $\bar{f} = 1$ (by definition).
2. $f = d + e \implies \frac{\partial f}{\partial d} = 1, \frac{\partial f}{\partial e} = 1. \bar{d} = \bar{f} \frac{\partial f}{\partial d} = 1 \cdot 1 = 1. \bar{e} = \bar{f} \frac{\partial f}{\partial e} = 1 \cdot 1 = 1.$
3. $d = \sqrt{c}, e = \cos(c)$. Node c receives gradient from its two children, d and e . $\bar{c} = \bar{d} \frac{\partial d}{\partial c} + \bar{e} \frac{\partial e}{\partial c} = 1 \cdot \frac{1}{2\sqrt{c}} + 1 \cdot (-\sin c) \approx \frac{1}{15.32} - 0.075 \approx 0.065 - 0.075 = -0.01.$

4. $c = a + b \implies \frac{\partial c}{\partial a} = 1, \frac{\partial c}{\partial b} = 1$. The gradient flows back from c to b : $\bar{b} = \bar{c} \frac{\partial c}{\partial b} = -0.01 \cdot 1 = -0.01$.
5. Node a is a parent to both b (via the exp operation) and c (via the addition). So, it receives gradient from both. $\bar{a} = \bar{b} \frac{\partial b}{\partial a} + \bar{c} \frac{\partial c}{\partial a} = \bar{b} \cdot (\exp a) + \bar{c} \cdot (1) \approx (-0.01) \cdot (54.6) + (-0.01) \cdot 1 \approx -0.546 - 0.01 = -0.556$.
6. **Final step (Gradient w.r.t. input x):** The gradient flows from a back to x .
 $\bar{x} = \bar{a} \frac{\partial a}{\partial x} = (-0.556) \cdot (2x) = -0.556 \cdot 4 = -2.224$.

Note on diagram values: The numerical values for the gradients in the text are derived step-by-step from the chain rule. These should be considered authoritative over the static values in the diagram if any discrepancies arise. The final result for $\frac{df}{dx}$ at $x = 2$ is approximately -2.224.

2.3 Application to Neural Networks

A neural network is just a very large, layered computation graph. Each layer performs a simple operation, and backpropagation works on this graph in exactly the same way. The output of layer i is the input to layer $i + 1$. A typical layer is:

$$\mathbf{a}_i = \sigma(W_i \mathbf{a}_{i-1} + \mathbf{b}_i)$$

where \mathbf{a}_{i-1} is the input from the previous layer, W_i, \mathbf{b}_i are the parameters (weights and biases) of the current layer, and σ is a non-linear activation function (like sigmoid, tanh, or ReLU).

The goal is to find the parameters $\theta = \{W_1, \mathbf{b}_1, \dots, W_L, \mathbf{b}_L\}$ that minimize a loss function, for example, the squared error:

$$L(\theta) = \|\mathbf{y}_{true} - \mathbf{a}_L\|^2$$

Backpropagation is the algorithm used to compute the gradient of this loss L with respect to every single parameter in the network $(\frac{\partial L}{\partial W_i}, \frac{\partial L}{\partial \mathbf{b}_i})$, allowing us to update them using gradient descent and train the model.





Figure 1: A cleaner view of a 2-layer neural network computation graph. The forward pass flows left-to-right through the data nodes (yellow circles). The parameters (blue rectangles) are updated via backpropagation (red dashed arrows), which computes their gradient with respect to the final loss.

3. Linearization: The Tangent Line Approximation

Often in math and engineering, we want to approximate a complex curve with a simple straight line, at least around a specific point of interest. This process is called **linearization**, and the gradient is the perfect tool for it. The linear approximation $L(\mathbf{x})$ of a function $f(\mathbf{x})$ at a point \mathbf{x}_0 is simply its tangent line (or plane, or hyperplane) at that point. It is given by the first two terms of the Taylor series:

$$L(\mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0)$$

This approximation is very accurate near \mathbf{x}_0 but naturally gets worse as you move further away. We use it to simplify hard problems into easy, linear ones, which is a cornerstone of many optimization algorithms.



3.1 Linearization Examples

Example 1: Find the linearization of $f(x) = \sqrt{x^2 + 9}$ around $x_0 = -4$.

1. Evaluate the function at the point: $f(-4) = \sqrt{(-4)^2 + 9} = \sqrt{16 + 9} = \sqrt{25} = 5$.
2. Find the derivative (gradient): $f'(x) = \frac{1}{2\sqrt{x^2+9}} \cdot (2x) = \frac{x}{\sqrt{x^2+9}}$.
3. Evaluate the derivative at the point: $f'(-4) = \frac{-4}{\sqrt{(-4)^2+9}} = \frac{-4}{5}$.
4. Assemble the linearization formula: $L(x) = f(x_0) + f'(x_0)(x - x_0)$.

$$L(x) = 5 - \frac{4}{5}(x - (-4)) = 5 - \frac{4}{5}(x + 4)$$

Example 2: Find the linearization of $f(x, y) = e^x \cos(y)$ around the point $(x_0, y_0) = (0, 0)$.

1. Evaluate function: $f(0, 0) = e^0 \cos(0) = 1 \cdot 1 = 1$.
2. Find the gradient vector: $\nabla f = \langle \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \rangle = \langle e^x \cos y, -e^x \sin y \rangle$.
3. Evaluate gradient at the point: $\nabla f(0, 0) = \langle e^0 \cos 0, -e^0 \sin 0 \rangle = \langle 1, 0 \rangle$.
4. Assemble linearization: $L(x, y) = f(0, 0) + \nabla f(0, 0) \cdot \langle x - 0, y - 0 \rangle$.

$$L(x, y) = 1 + \langle 1, 0 \rangle \cdot \langle x, y \rangle = 1 + (1 \cdot x + 0 \cdot y) = 1 + x$$

4. Practice Problems

Problem 1

We define

$$g(\mathbf{z}, \nu) := \log p(\mathbf{x}, \mathbf{z}) - \log q(\mathbf{z}, \nu)$$

$$\mathbf{z} := t(\epsilon, \nu)$$

for differentiable functions p, q, t , and $\mathbf{x} \in \mathbb{R}^D, \mathbf{z} \in \mathbb{R}^E, \nu \in \mathbb{R}^F, \epsilon \in \mathbb{R}^G$. By using the chain rule, compute the total derivative

$$\frac{d}{d\nu} g(\mathbf{z}, \nu).$$

Problem 2

Compute the derivatives $\frac{df}{dx}$ of the following functions by using the chain rule. Provide the dimensions of every single partial derivative. Describe your steps in detail.

a.

$$f(z) = \log(1 + z), \quad z = \mathbf{x}^\top \mathbf{x}, \quad \mathbf{x} \in \mathbb{R}^D$$

b.

$$\mathbf{f}(\mathbf{z}) = \sin(\mathbf{z}), \quad \mathbf{z} = A\mathbf{x} + \mathbf{b}, \quad A \in \mathbb{R}^{E \times D}, \mathbf{x} \in \mathbb{R}^D, \mathbf{b} \in \mathbb{R}^E$$

where $\sin(\cdot)$ is applied to every element of \mathbf{z} , and the output \mathbf{f} is also in \mathbb{R}^E .

5. Solutions to Practice Problems

Solution to Problem 1 - A More Intuitive Approach

Goal: Find the total effect that changing the parameter vector ν has on the final scalar value g .

Strategy: Think of ν as a set of control knobs. Turning these knobs has two distinct effects on g , and we must account for both:

1. **Direct Path:** ν appears directly inside the function g (within the $\log q$ term). Changing ν has an immediate, direct impact on g .
2. **Indirect Path:** ν is also used to compute the intermediate variable \mathbf{z} (since $\mathbf{z} = t(\epsilon, \nu)$). Changing ν changes \mathbf{z} , and this new value of \mathbf{z} then propagates forward to change the value of g .

To find the total derivative (the total sensitivity), we must calculate the gradient from each path and add them together. This is the essence of the multivariate chain rule.

Step-by-Step Execution: The formula for the total derivative is a sum over all paths of influence:

$$\frac{dg}{d\nu} = \underbrace{\frac{\partial g}{\partial \nu}}_{\text{Path 1: Direct influence of } \nu \text{ on } g} + \underbrace{\frac{\partial g}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \nu}}_{\text{Path 2: Indirect, via } \mathbf{z}}$$

- **Calculate Path 1 (Direct Effect):** We find the partial derivative of g with

respect to ν , treating \mathbf{z} as a constant for this step.

$$\frac{\partial g}{\partial \nu} = \frac{\partial}{\partial \nu} (\log p(\mathbf{x}, \mathbf{z}) - \log q(\mathbf{z}, \nu)) = -\frac{\partial \log q(\mathbf{z}, \nu)}{\partial \nu}$$

- **Calculate Path 2 (Indirect Effect):** This path itself has two links in the chain.

– First, how sensitive is g to changes in \mathbf{z} ?

$$\frac{\partial g}{\partial \mathbf{z}} = \frac{\partial \log p(\mathbf{x}, \mathbf{z})}{\partial \mathbf{z}} - \frac{\partial \log q(\mathbf{z}, \nu)}{\partial \mathbf{z}}$$

– Second, how sensitive is \mathbf{z} to changes in ν ?

$$\frac{\partial \mathbf{z}}{\partial \nu} = \frac{\partial t(\epsilon, \nu)}{\partial \nu}$$

Final Answer: Now we combine everything to get the total derivative.

$$\frac{dg}{d\nu} = -\frac{\partial \log q(\mathbf{z}, \nu)}{\partial \nu} + \left(\frac{\partial \log p(\mathbf{x}, \mathbf{z})}{\partial \mathbf{z}} - \frac{\partial \log q(\mathbf{z}, \nu)}{\partial \mathbf{z}} \right) \frac{\partial t(\epsilon, \nu)}{\partial \nu}$$

Solution to Problem 2a

Goal: Find the gradient of the scalar output f with respect to the vector input \mathbf{x} .

Strategy: This is a two-step computation graph: $\mathbf{x} \rightarrow z \rightarrow f$. The vector \mathbf{x} first gets converted to a scalar z , which then produces the final scalar f . The chain rule lets us multiply the sensitivities (derivatives) of each step.

$$\frac{df}{d\mathbf{x}} = (\text{sensitivity of } f \text{ to } z) \times (\text{sensitivity of } z \text{ to } \mathbf{x})$$

$$\frac{df}{d\mathbf{x}} = \frac{df}{dz} \frac{\partial z}{\partial \mathbf{x}}$$

Step-by-Step Execution:

1. **First Link ($\frac{df}{dz}$):** Sensitivity of the final output f to its immediate input z .

This is a standard scalar derivative.

$$\frac{df}{dz} = \frac{d}{dz} \log(1 + z) = \frac{1}{1 + z}$$

Dimension: Scalar (1×1).

2. **Second Link ($\frac{\partial z}{\partial \mathbf{x}}$):** Sensitivity of the intermediate scalar z to the input vector \mathbf{x} . This is the derivative of a scalar by a vector, which results in a

row vector. Let's derive it from scratch: $z = \mathbf{x}^\top \mathbf{x} = \sum_{i=1}^D x_i^2$. The k -th component of the gradient is $\frac{\partial z}{\partial x_k} = 2x_k$. Arranging these into a row vector gives $[2x_1, 2x_2, \dots, 2x_D] = 2\mathbf{x}^\top$.

$$\frac{\partial z}{\partial \mathbf{x}} = 2\mathbf{x}^\top$$

Dimension: Row vector ($1 \times D$).

Final Answer: We multiply the results (scalar times row vector). The dimensions $(1 \times 1) \cdot (1 \times D)$ yield a $(1 \times D)$ row vector, as expected for the gradient of a scalar function with respect to a column vector.

$$\frac{df}{d\mathbf{x}} = \left(\frac{1}{1+z} \right) (2\mathbf{x}^\top) = \frac{2\mathbf{x}^\top}{1 + \mathbf{x}^\top \mathbf{x}}$$

Solution to Problem 2b

Goal: Find the derivative of the vector function \mathbf{f} with respect to the vector input \mathbf{x} . When the input and output are both vectors, the derivative is a matrix called the **Jacobian**.

Strategy: The computation graph is $\mathbf{x} \rightarrow \mathbf{z} \rightarrow \mathbf{f}$. The chain rule for Jacobians works like matrix multiplication of the Jacobians from each step.

$$\underbrace{\frac{\partial \mathbf{f}}{\partial \mathbf{x}}}_{\text{Jacobian of } \mathbf{f} \text{ wrt } \mathbf{x}} = \underbrace{\frac{\partial \mathbf{f}}{\partial \mathbf{z}}}_{\text{Jacobian of } \mathbf{f} \text{ wrt } \mathbf{z}} \cdot \underbrace{\frac{\partial \mathbf{z}}{\partial \mathbf{x}}}_{\text{Jacobian of } \mathbf{z} \text{ wrt } \mathbf{x}}$$

Step-by-Step Execution:

1. **First Link** ($\frac{\partial \mathbf{f}}{\partial \mathbf{z}}$): This is the Jacobian of \mathbf{f} with respect to \mathbf{z} . The key insight here is that the $\sin(\cdot)$ function is applied **element-wise**: $f_i = \sin(z_i)$. This means that f_i depends *only* on z_i , and not on any other z_j where $j \neq i$. Therefore, the partial derivative $\frac{\partial f_i}{\partial z_j}$ is 0 if $i \neq j$. The Jacobian matrix is zero everywhere except for the main diagonal.

$$(\frac{\partial \mathbf{f}}{\partial \mathbf{z}})_{ij} = \frac{\partial f_i}{\partial z_j} = \begin{cases} \cos(z_i) & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

This forms a diagonal matrix:

$$\frac{\partial \mathbf{f}}{\partial \mathbf{z}} = \text{diag}(\cos(z_1), \dots, \cos(z_E)) = \text{diag}(\cos(\mathbf{z}))$$

Dimension: Square diagonal matrix ($E \times E$).

2. **Second Link ($\frac{\partial \mathbf{z}}{\partial \mathbf{x}}$)**: This is the Jacobian of \mathbf{z} with respect to \mathbf{x} . For the linear transformation $\mathbf{z} = A\mathbf{x} + \mathbf{b}$, the derivative is simply the matrix A .

$$\frac{\partial \mathbf{z}}{\partial \mathbf{x}} = A$$

Dimension: The matrix A itself ($E \times D$).

Final Answer: We multiply the two matrices. The order is crucial. The dimensions $(E \times E) \cdot (E \times D)$ yield a final $(E \times D)$ Jacobian matrix.

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \text{diag}(\cos(\mathbf{z})) \cdot A = \text{diag}(\cos(A\mathbf{x} + \mathbf{b})) \cdot A$$

Conclusion

We have journeyed from a simple analogy of playing darts to the formal mechanics of how machine learning models truly learn. The two key takeaways should be:

1. **Backpropagation is just the chain rule applied to a computation graph.**
By breaking down complex functions into simple steps, we can efficiently compute the gradient of every parameter by passing gradient information backward through the graph.
2. **Linearization is approximation using the gradient.** It allows us to replace a complex function with its tangent line at a point, simplifying analysis and forming the basis for powerful optimization methods.

Mastering these two concepts provides a solid foundation for understanding and developing advanced machine learning and deep learning models.